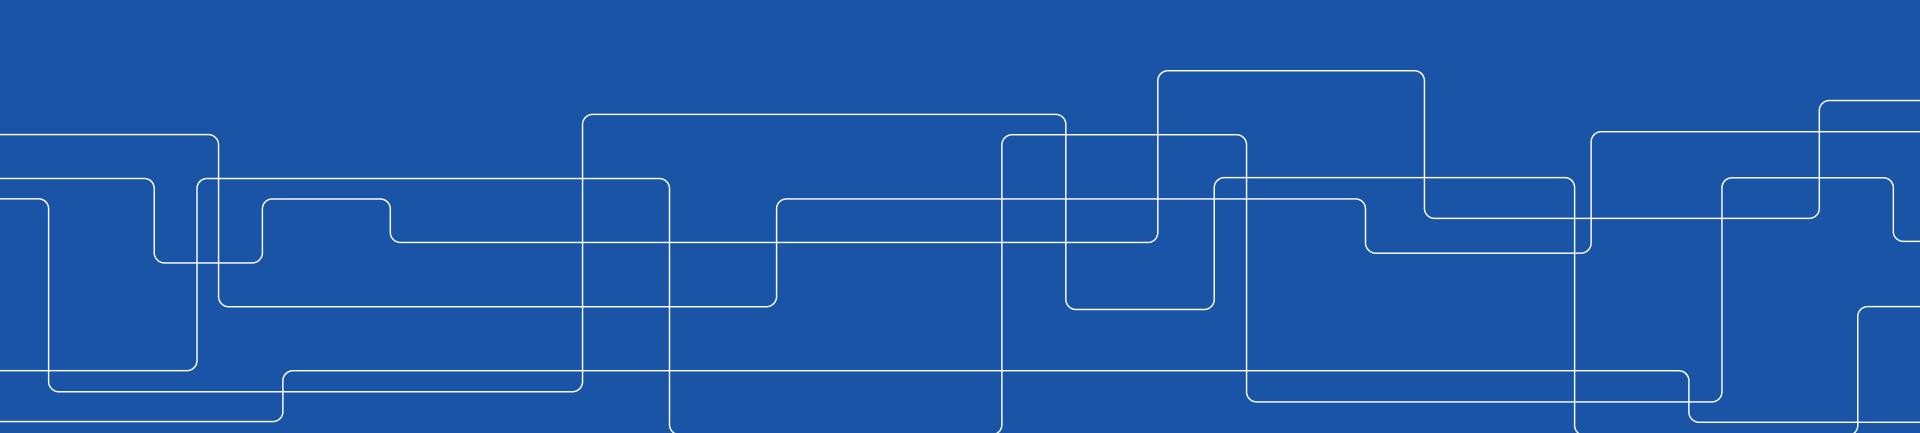




Remote Invocation

Vladimir Vlassov and Johan Montelius





Middleware

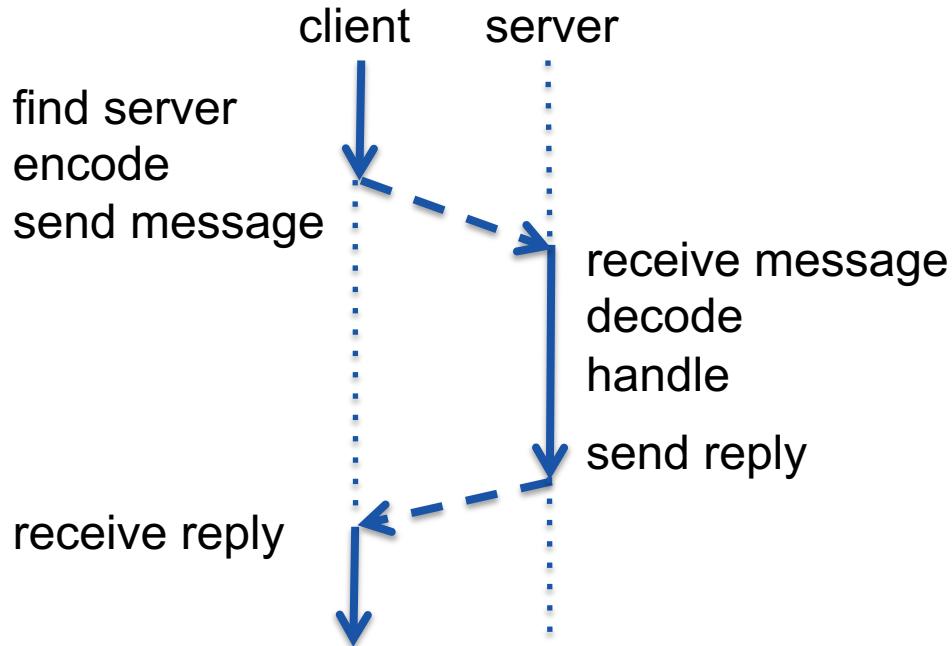
Application layer

Remote invocation / indirect communication

Socket layer

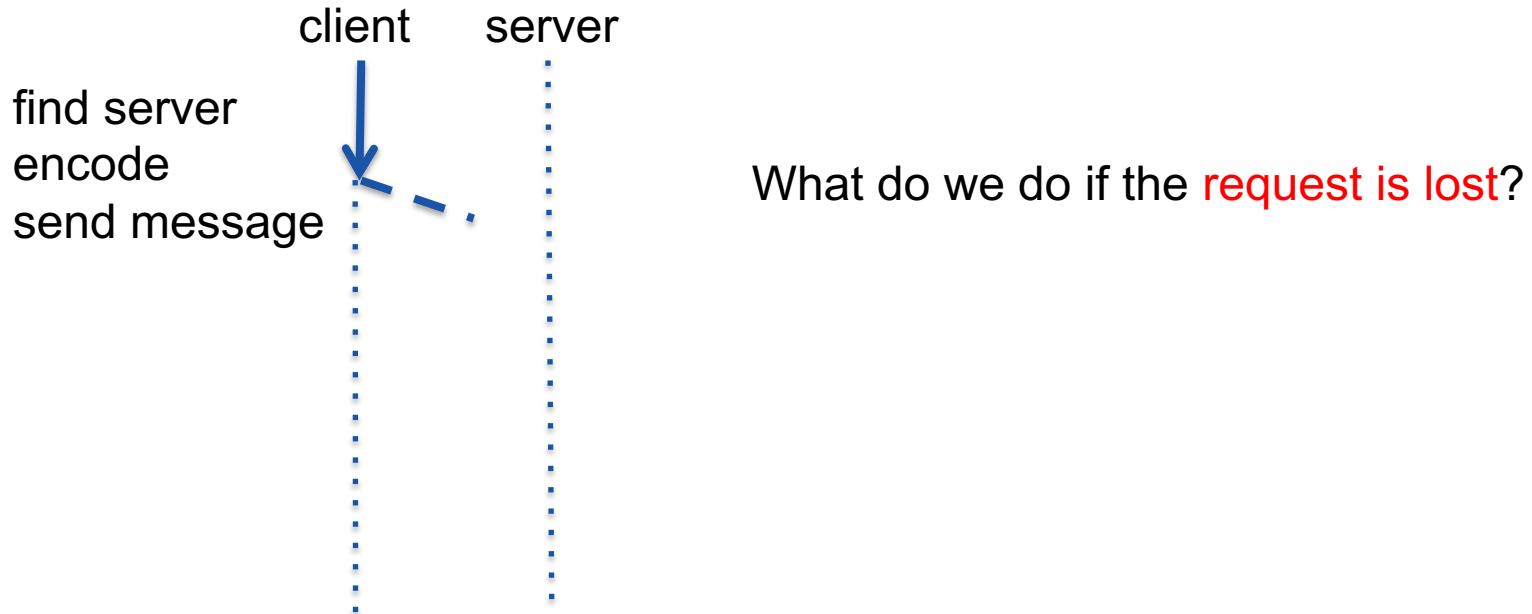
Network layer

Request / Reply

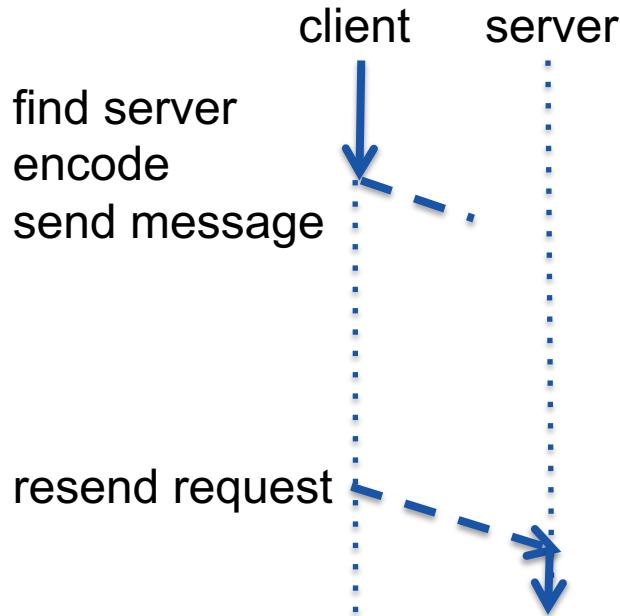


- identify and locate the server
- encode/decode the message
- send a reply to the right client
- attach a reply to the request

Lost request

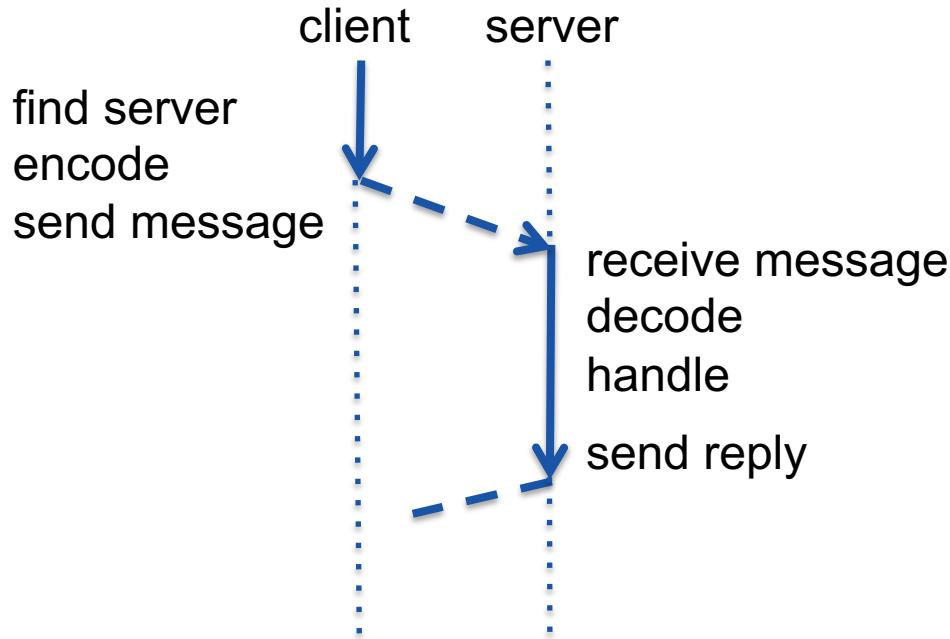


Resend request



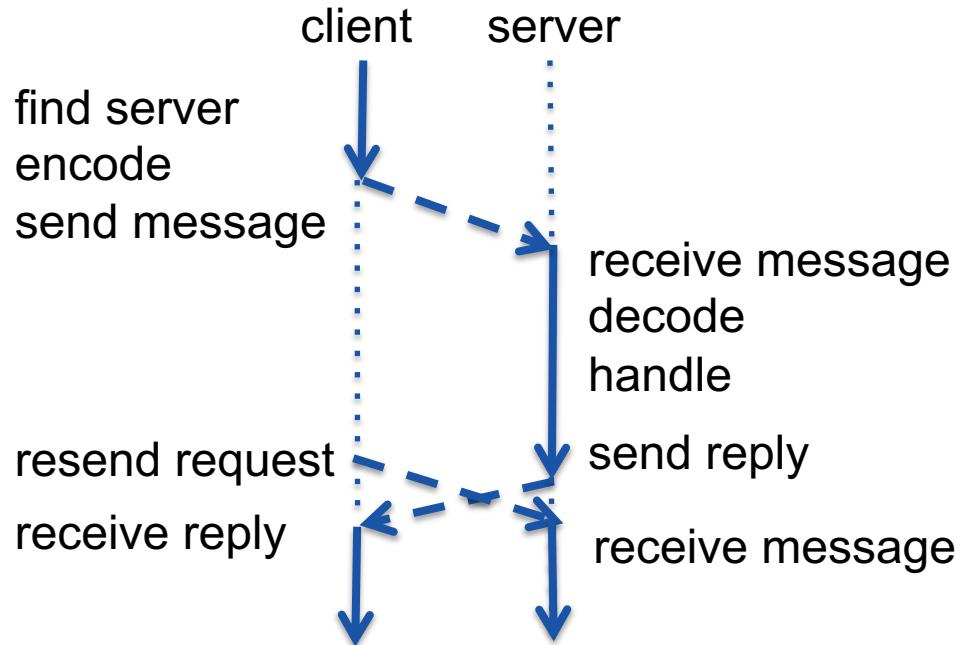
- need to detect that message is potentially lost
- wait for a **timeout** (how long) or error from the underlying layer
- **resend** the request
- simple, problem solved

Lost reply



- the client will wait for a **timeout** and **re-send the request**
- not a problem

Problem



- a problem
- the server **might need a history of all previous request**
- *might need*



Idempotent operations

- add 100 euros to my account
- what is the status of my account
- Sweden scored yet another goal!
- The standing is now 2-1!



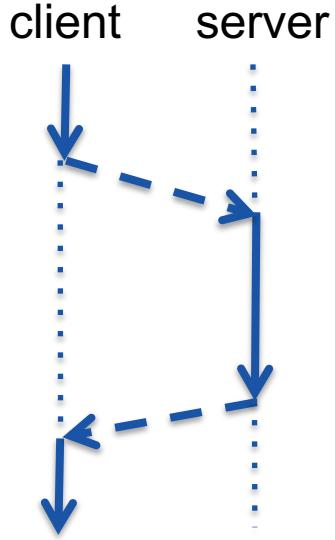
History

If operations are **not idempotent**, the server must ensure that the same request is **not executed twice**.

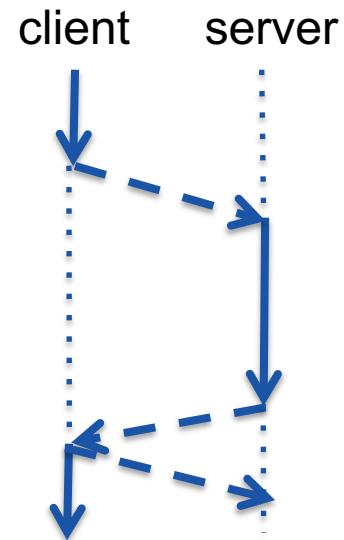
Keep a **history of all requests and replies**. The same reply can be sent without re-executing if a request is resent.

For how long do you keep the history?

Request-Reply-Acknowledge



Request-Reply (RR)



Request-Reply-Acknowledge (RRA)



At-most-once or At-least-once

How about this:

If an operation **succeeds**, then...

At-most-once: the request has been **executed once**.

Implemented using a history or simply not resending requests.

At-least-once: the request has been **executed at least once**.

No need for a history; simply resend requests until a reply is received.



At most or At least

How about **errors**:

Even if we do resend messages, we will have to give up at some time.

If an operation **fails/is lost**, then...

at-most-once:

- the request might have been executed at most once

at-least-once:

- the request might have been executed at least once

At most or At least

Pros and cons:

- *At-most-once without re-sending requests:*
Simple to implement, not fault-tolerant
- *At-most-once with history:*
Expensive to implement, fault-tolerant
- *At-least-once:*
Simple to implement, fault-tolerant

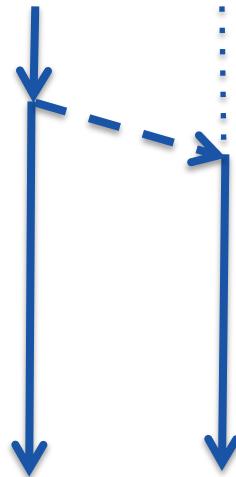
Can you live with at-least-once semantics?



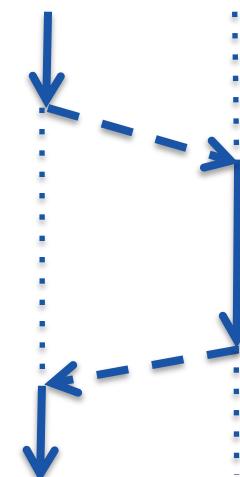
UDP or TCP

Should we implement a request-reply protocol over UDP or TCP?

Synchronous or Asynchronous

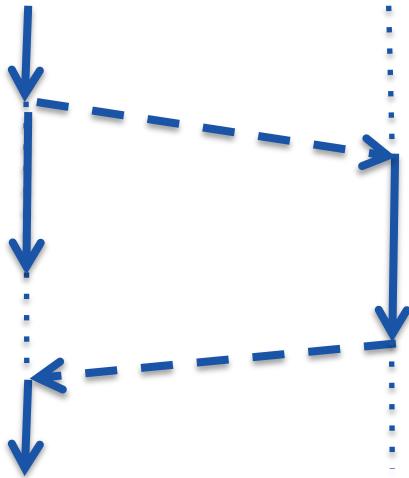


Asynchronous



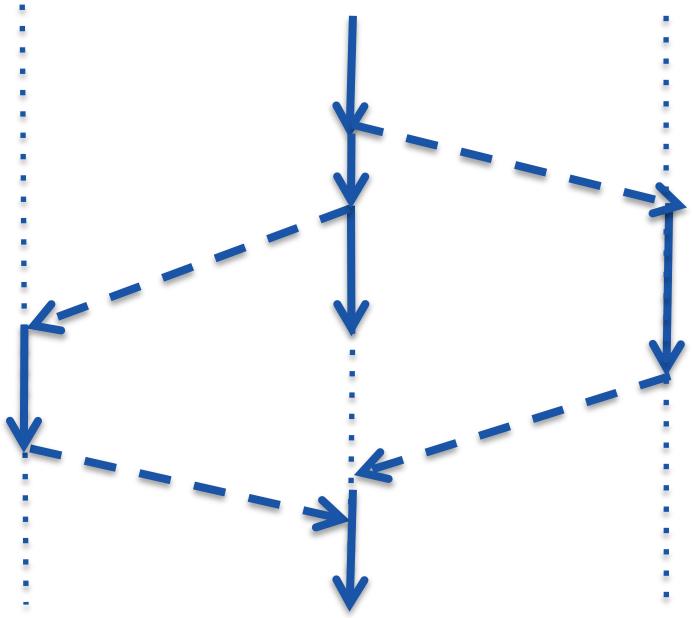
Synchronous

RR over Asynchronous



- send request
- continue to execute
- suspend if not arrived
- read reply

Hide the latency





HTTP

A request-reply protocol described in RFC 2616.

Request = Request-Line *(header CRLF) CRLF [message-body]

Request-Line = Method SP Request-URI SP HTTP-Version CRLF

GET /index.html HTTP/1.1\r\n foo 42 \r\n\r\nHello

HTTP methods

- **GET**: request a resource, *should be idempotent*
- **HEAD**: request only header information
- **POST**: upload information to a resource, included in the body, status of the server could change
- **PUT**: add or replace a resource, idempotent
- **DELETE**: add or replace content, idempotent

Wireshark

Frame 74: 699 bytes on wire (5592 bits), 699 bytes captured (5592 bits) on interface 0

Ethernet II, Src: AsustekC_93:c6:da (00:1e:8c:93:c6:da), Dst: All-HSRP-routers_d4 (00:00:0c:07:ac:d4)

Internet Protocol Version 4, Src: 130.237.215.140 (130.237.215.140), Dst: 130.237.28.40 (130.237.28.40)

Transmission Control Protocol, Src Port: 53960 (53960), Dst Port: 80 (80), Seq: 1, Ack: 1, Len: 645

Hypertext Transfer Protocol

GET / HTTP/1.1\r\n

Host: www.kth.se\r\n

User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:40.0) Gecko/20100101 Firefox/40.0\r\n

Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8\r\n

Accept-Language: en-US,en;q=0.5\r\n

Accept-Encoding: gzip, deflate\r\n

[truncated]Cookie: __utma=154244322.999183788.1409574123.1430294703.1441199888.6; modalVisitorPoll=participate#1; csrfToken=rw\r\n

Connection: keep-alive\r\n

[Full request URI: http://www.kth.se/]

[HTTP request 1/1]

0050	6b 74 68 2e 73 65 0d 0a	55 73 65 72 2d 41 67 65	kth.se.. User-Age
0060	6e 74 3a 20 4d 6f 7a 69	6c 6c 61 2f 35 2e 30 20	nt: Mozilla/5.0
0070	28 58 31 31 32 20 55 62	75 6e 74 75 3b 20 4c 69	(X11; Ubuntu; Li
0080	6e 75 78 20 78 38 36 5f	36 34 3b 20 72 76 3a 34	nux x86_64; rv:4
0090	30 2e 30 29 20 47 65 63	6b 6f 2f 32 30 31 30 30	0.0) Gecko/20100
00a0	31 30 31 20 46 69 72 65	66 6f 78 2f 34 30 2e 30	101 Firefox/40.0
00b0	0d 0a 41 63 63 65 70 74	3a 20 74 65 78 74 2f 68	Accept: text/h

HTTP User-Agent header (http.user...) Packets: 1017 · Displayed: 1017 (100,0%) · Dropped: 0 (0,0%) Profile: Default



HTTP GET

GET / HTTP/1.1

Host: www.kth.se

User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:40.0) Gecko/20100101

Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8

Accept-Language: en-US,en;q=0.5

Accept-Encoding: gzip, deflate

Cookie:

Connection: keep-alive



HTTP Response

HTTP/1.1 200 OK

Date: Tue, 08 Sep 2015 10:37:49 GMT

Server: Apache/2.2.15 (Red Hat)

X-UA-Compatible: IE=edge

Set-Cookie: JSESSIONID=CDC76A3;Path=/; Secure; HttpOnly

Content-Language: sv-SE

Content-Length: 59507

Connection: close

Content-Type: text/html;charset=UTF-8

<!DOCTYPE html>

<html lang="sv">

<title>KTH | Valkommen till KTH</title>



The web

On the web, the resource is often an HTML document that is presented in a browser.

HTTP could be used as a general-purpose request-reply protocol.



REST and SOAP

Request-reply protocols for Web services:

- **REST (Representational State Transfer)**
 - content described in XML, JSON, . . .
 - lightweight,
- **SOAP (Simple Object Access Protocol)**
 - over HTTP, SMTP . . .
 - content described in SOAP/XML
 - standardized, heavyweight



HTTP over TCP

HTTP over TCP - a good idea?



Masking a request-reply

Could we use a regular program construct to **hide** the fact that we do a **request-reply**?



Masking a request-reply

Could we use a regular program construct to **hide** the fact that we do a **request-reply**?

- **RPC**: Remote Procedure Call
- **RMI**: Remote Method Invocation

Motivation for RPC and RMI

Message passing is convenient for consumers-producers (filters) and P2P, but it is somewhat low-level for client-server applications.

- Client/server interactions are based on a request/response protocol;
- Client requests are typically mapped to procedures on the server;
- A client waits for a response from the server.

Need for more convenient (easier to use) communication mechanisms for developing client/server applications



Motivation for RPC and RMI

Remote Procedure Call (RPC) and rendezvous

- Procedure interface; message passing implementation

Remote Method Invocation (RMI)

- RMI is an object-oriented analog of RPC

RPC, rendezvous, and RMI are implemented on top of message passing.



Procedure calls

What is a procedure call:

- find the procedure
- give the procedure access to arguments
- pass control to the procedure
- collect the reply if any
- continue execution

How do we turn this into **a tool for distributed programming?**



Operational semantics

```
int x, n;  
n = 5;  
proc(n);  
x = n;
```

```
int x, arr[3];  
arr[0] = 5;  
proc(arr);  
x = arr[0];
```



Call by value/reference

Call by value

- A procedure is given a copy of the datum

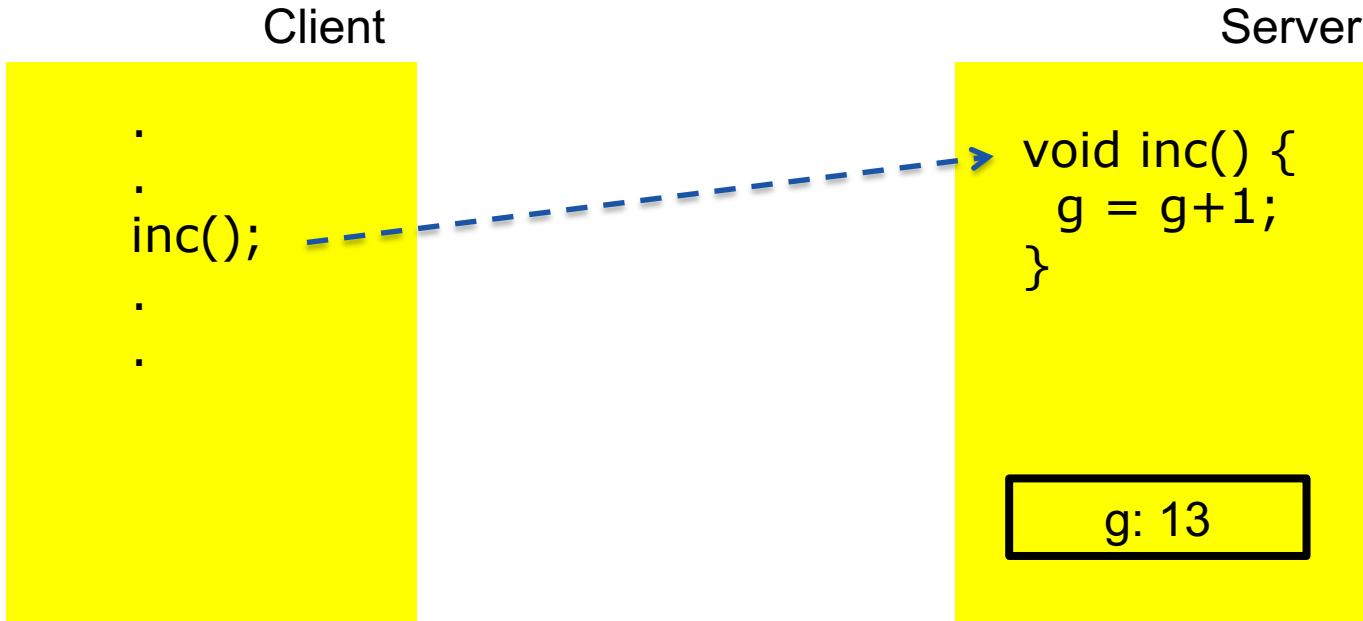
Call by reference

- A procedure is given a reference to the datum

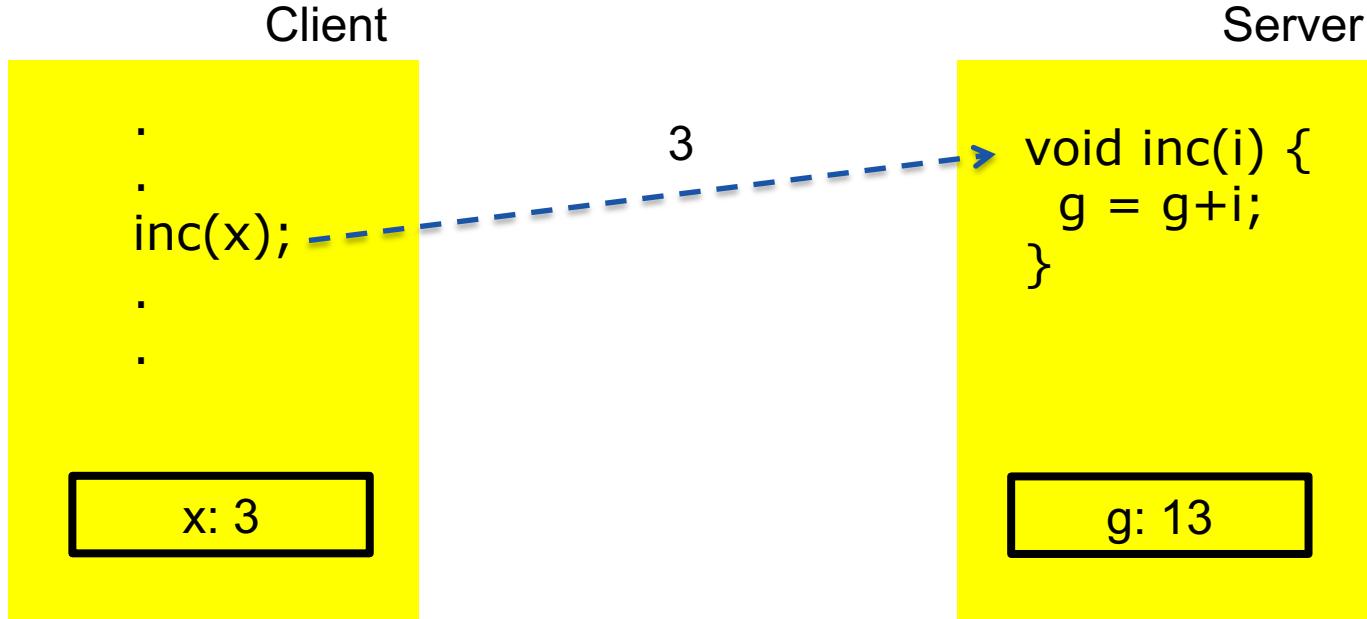
What if the datum is a reference, and we pass a copy of the datum?

Why is this important?

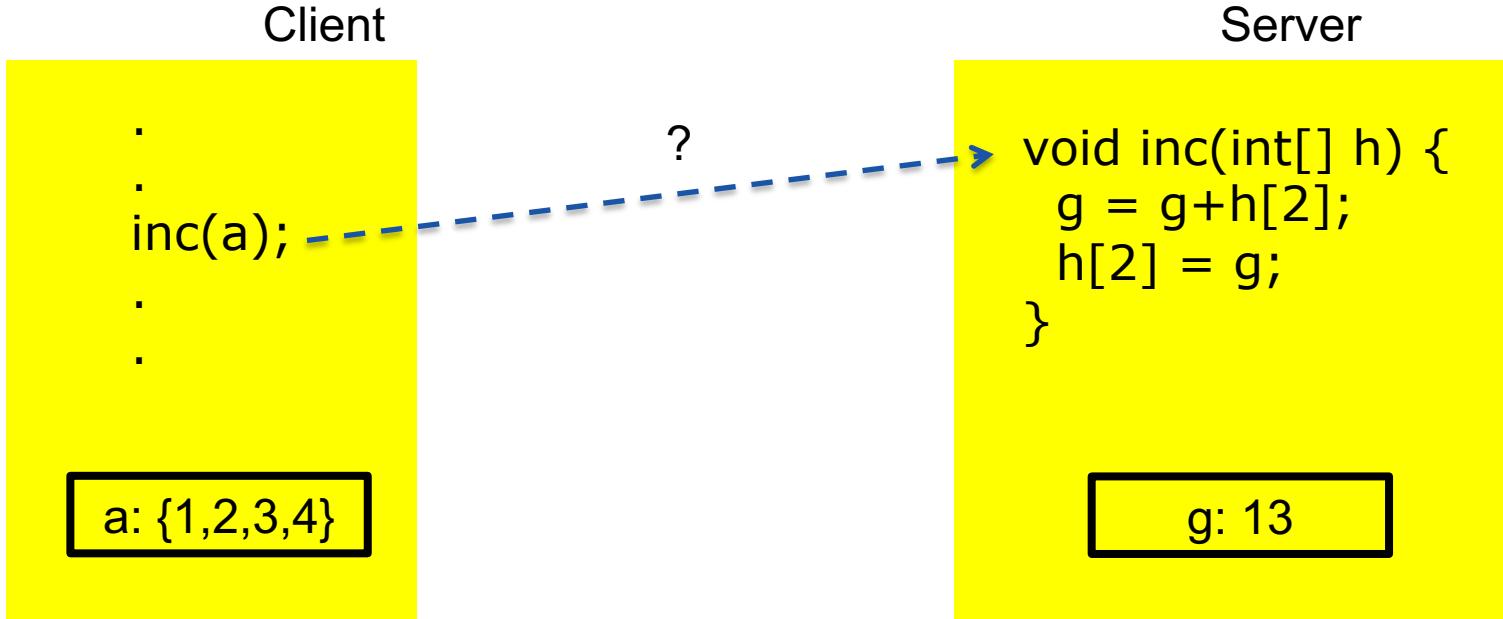
RPC: Remote Procedure Call



RPC: Remote Procedure Call



RPC: Remote Procedure Call

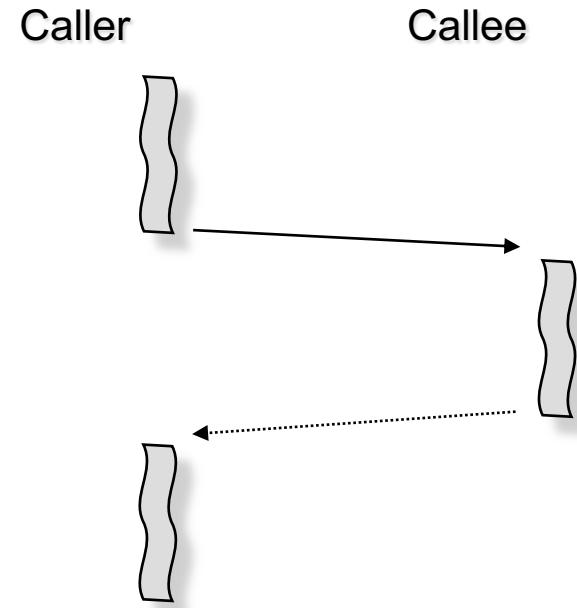


RPC: Remote Procedure Call

RPC is a mechanism that allows a program running on one computer (VM) to cause a procedure to be executed on another computer (VM) without the programmer needing to code for this explicitly.

Two processes are involved:

- **Caller (RPC client)** is a **calling process** that initiates an RPC to a server.
- **Callee (RPC server)** is a **called process** that accepts the call.

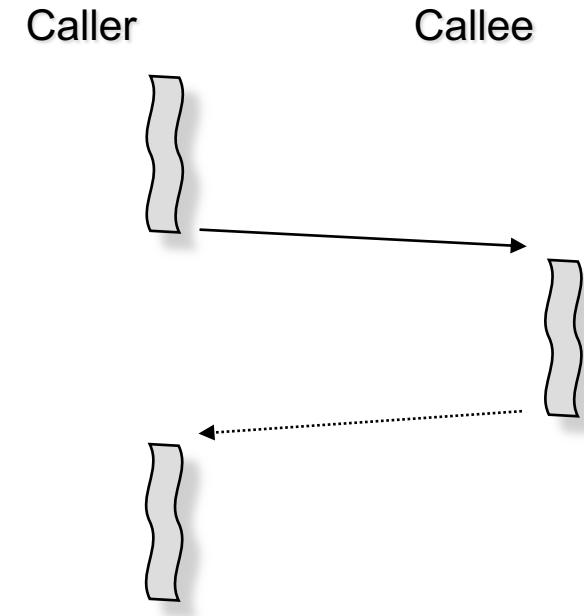


RPC: Remote Procedure Call (cont'd)

Each RPC is executed in a **separate process (thread)** on the server side

An RPC is a synchronous operation.

- The caller is suspended until the results of the remote procedure are returned.
- Like a regular or local procedure call.
- Guess why?



Identifying a Remote Procedure

Each RPC procedure is uniquely identified by

- A program number
 - identifies a group of related remote procedures
- A version number
- A procedure number

An RPC call message has three unsigned fields:

- Remote program number
- Remote program version number
- Remote procedure number

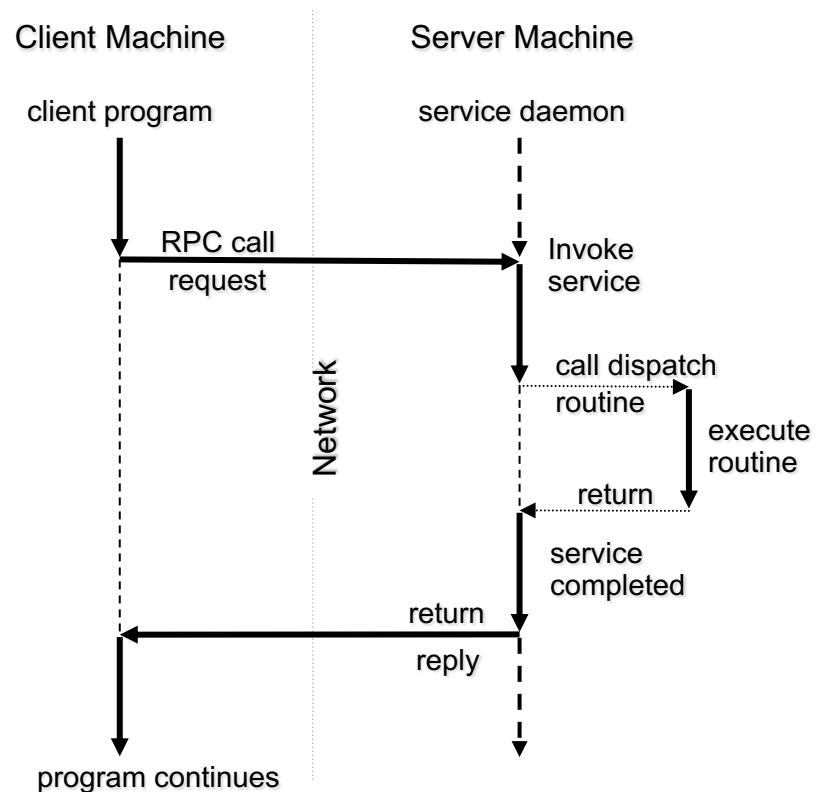
The three fields uniquely identify the procedure to be called.

Executing RPC

On each RPC, the server starts a ***new process*** to execute the call.

- The new process terminates when the procedure returns and results are sent to the caller.
- Calls from the same caller and calls from different callers are serviced by ***different concurrent processes*** on the server.

Concurrent invocations might interfere when accessing shared objects – might need ***synchronization***.



An RPC Syntax

- Modules (Servers)

```
module mname
    interface, i.e. headers of exported operations;
body
    variable declarations;
    initialization code;
    procedures for exported operations;
    local procedures and processes;
end mname
```

- Exported operation (method of a remote interface)

```
op opname(formal identifiers) [returns result]
```

- Procedure – operation implementation

```
proc opname(formal identifiers) returns result identifier
    declarations of local variables;
    statements
end
```

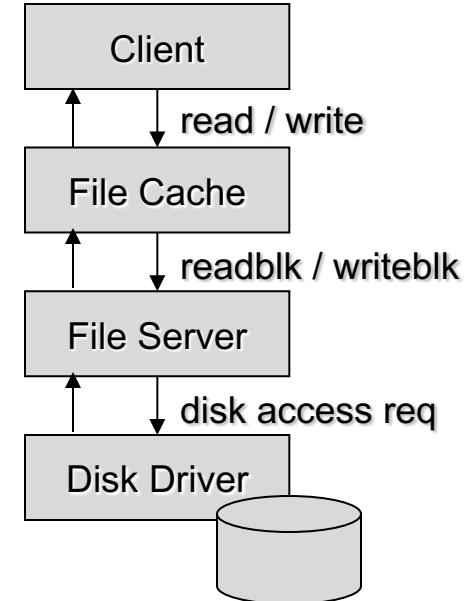
- Client makes a remote call to a module (server):

```
call mname.opname(arguments)
```

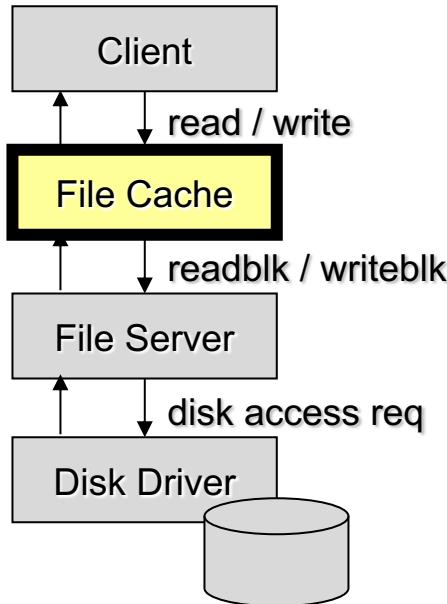
Example: A Distributed File System Using RPC

Modules (servers):

- **FileCache**
 - A write-back allocate-on-write cache of file blocks
 - Exports **read** and **write** operations
 - On a miss, calls remote **FileServer**
- **FileServer**
 - Provides access to file blocks stored on a disk
 - Exports **readblk** and **writeblk**
 - Uses the local **DiskDriver** process to access the disk.



The File Cache



```

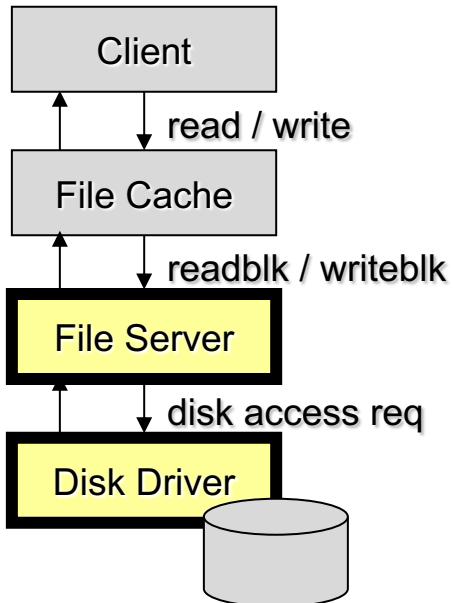
module FileCache  # located on each diskless workstation
    op read(int count; result char buffer[*]);
    op write(int count; char buffer[]);
body
cache of file blocks;
variables to record file descriptor information;
semaphores for synchronization of cache access (if needed);

proc read(count,buffer) {
    if (needed data is not in cache) {
        select cache block to use;
        if (need to write out the cache block)
            FileServer.writeblk(...);
            FileServer.readblk(...);
    }
    buffer = appropriate count bytes from cache block;
}

proc write(count,buffer) {
    if (appropriate block not in cache) {
        select cache block to use;
        if (need to write out the cache block)
            FileServer.writeblk(...);
    }
    cache block = count bytes from buffer;
}
end FileCache

```

File Server and Disk Driver



```
module FileServer    # located on a file server
  op readblk(int fileid, offset; result char blk[1024]);
  op writeblk(int fileid, offset; char blk[1024]);
body
  cache of disk blocks;
  queue of pending disk access requests;
  semaphores to synchronize access to the cache and queue;
  # N.B. synchronization code not shown below

  proc readblk(fileid, offset, blk) {
    if (needed block not in the cache) {
      store read request in disk queue;
      wait for read operation to be processed;
    }
    blk = appropriate disk block;
  }

  proc writeblk(fileid, offset, blk) {
    select block from cache;
    if (need to write out the selected block) {
      store write request in disk queue;
      wait for block to be written to disk;
    }
    cache block = blk;
  }

  process DiskDriver {
    while (true) {
      wait for a disk access request;
      start a disk operation; wait for interrupt;
      awaken process waiting for this request to complete;
    }
  }
end FileServer
```

Open Network Computing (ONC) RPC (SunRPC)

- targeting intranet, file servers, etc
- **at-least-once** call semantics
- procedures described in **Interface Definition Language (IDL)**
- XDR (eXternal Data Representation) specifies message structure
- used UDP as transport protocol (TCP also available)



Java RMI (Remote Method Invocation)

- Similar to RPC but:
 - we now **invoke methods of remote objects**
 - **at-most-once** semantics
- Objects can be passed as arguments; how should this be done?
 - **by value**
 - **by reference**



Java RMI

We can do either:

A *remote object* is passed as a reference (*by reference*), i.e., it remains at the original place it was created.

A *serializable object* is passed as a copy (*by value*), i.e., the object is duplicated.



Finding the procedure/object

How do we locate a remote procedure/object/process?

Network address that specifies the location or...

A known “binder” process that keeps track of registered resources.

Remote Method Invocation (RMI)

Remote method invocation (RMI) is a mechanism to invoke a method on a remote object, i.e., an object in another computer or virtual machine.

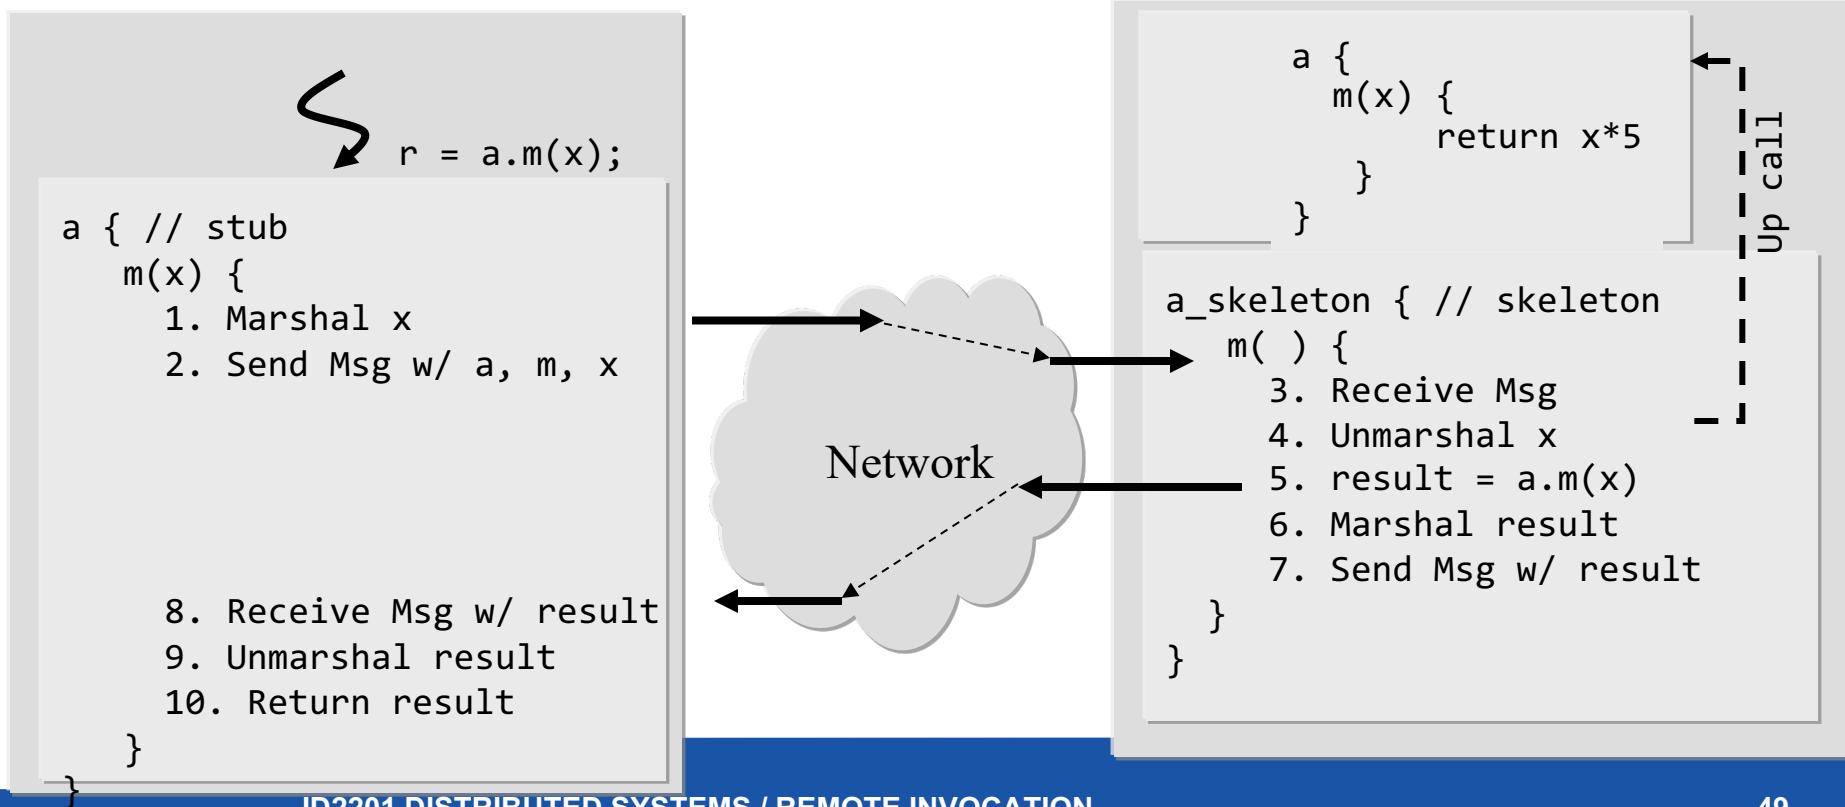
RMI is the object-oriented analog of RPC in a distributed OO environment, e.g., OMG CORBA, Java RMI, .NET

- RPC allows calling procedures over a network
- RMI invokes objects' methods over a network

Location transparency: invoke a method on a stub like on a local object

Location awareness: the stub makes remote calls across a network and returns results via stack

Remote Method Invocation





Locating Objects

How does a caller get a reference to a remote object, i.e., stub?

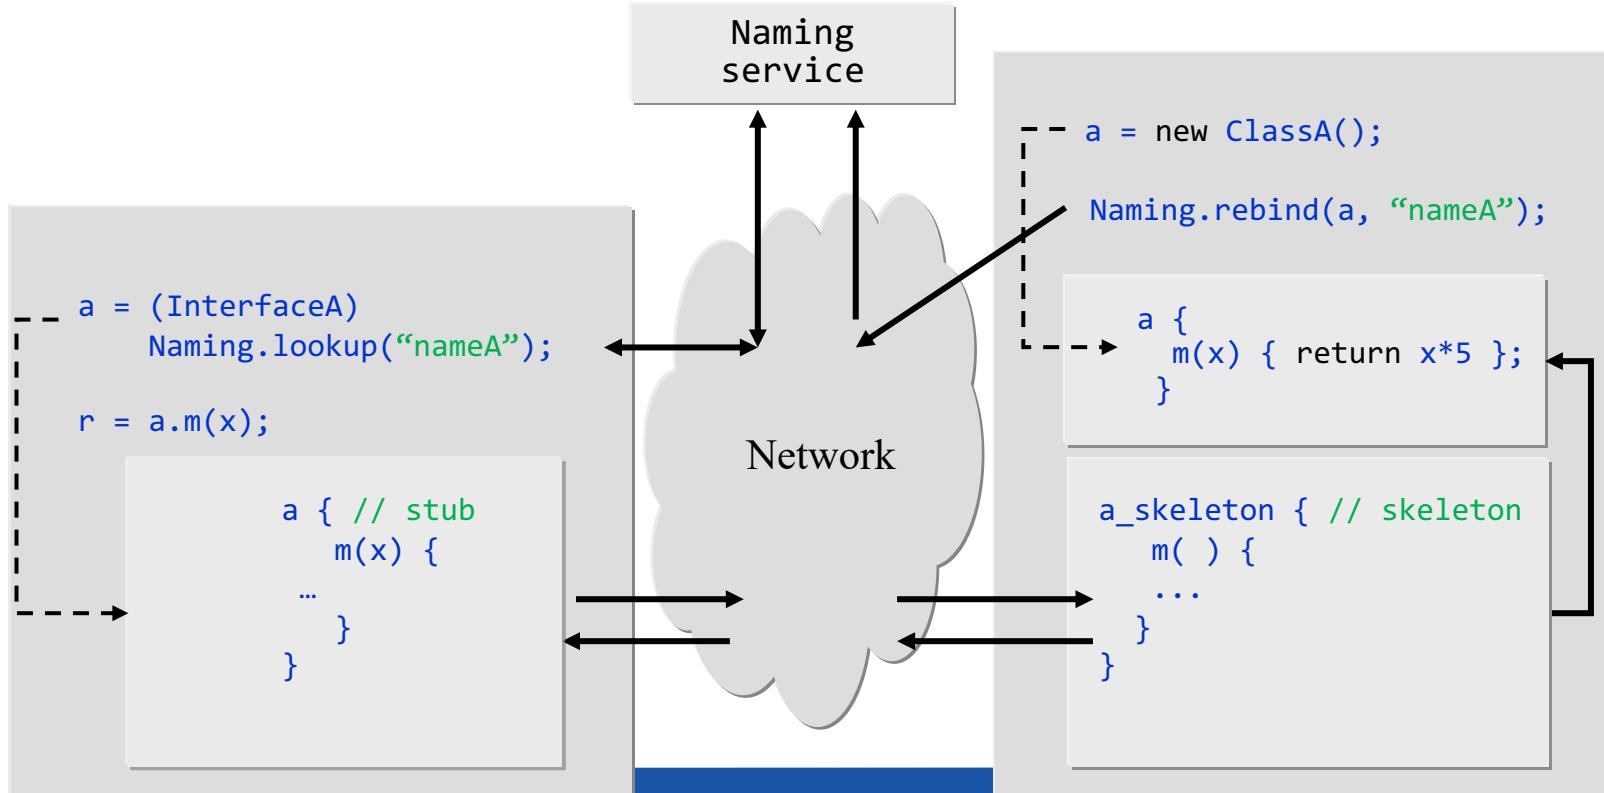
One approach is to use a **Naming Service**:

- Associate a unique name with an object.
- Bind the name to the object at the Naming Service.
 - The record typically includes name, class name, object reference (i.e., location information), and other information to create a stub.
- The client looks up the object by name in the Naming Service.

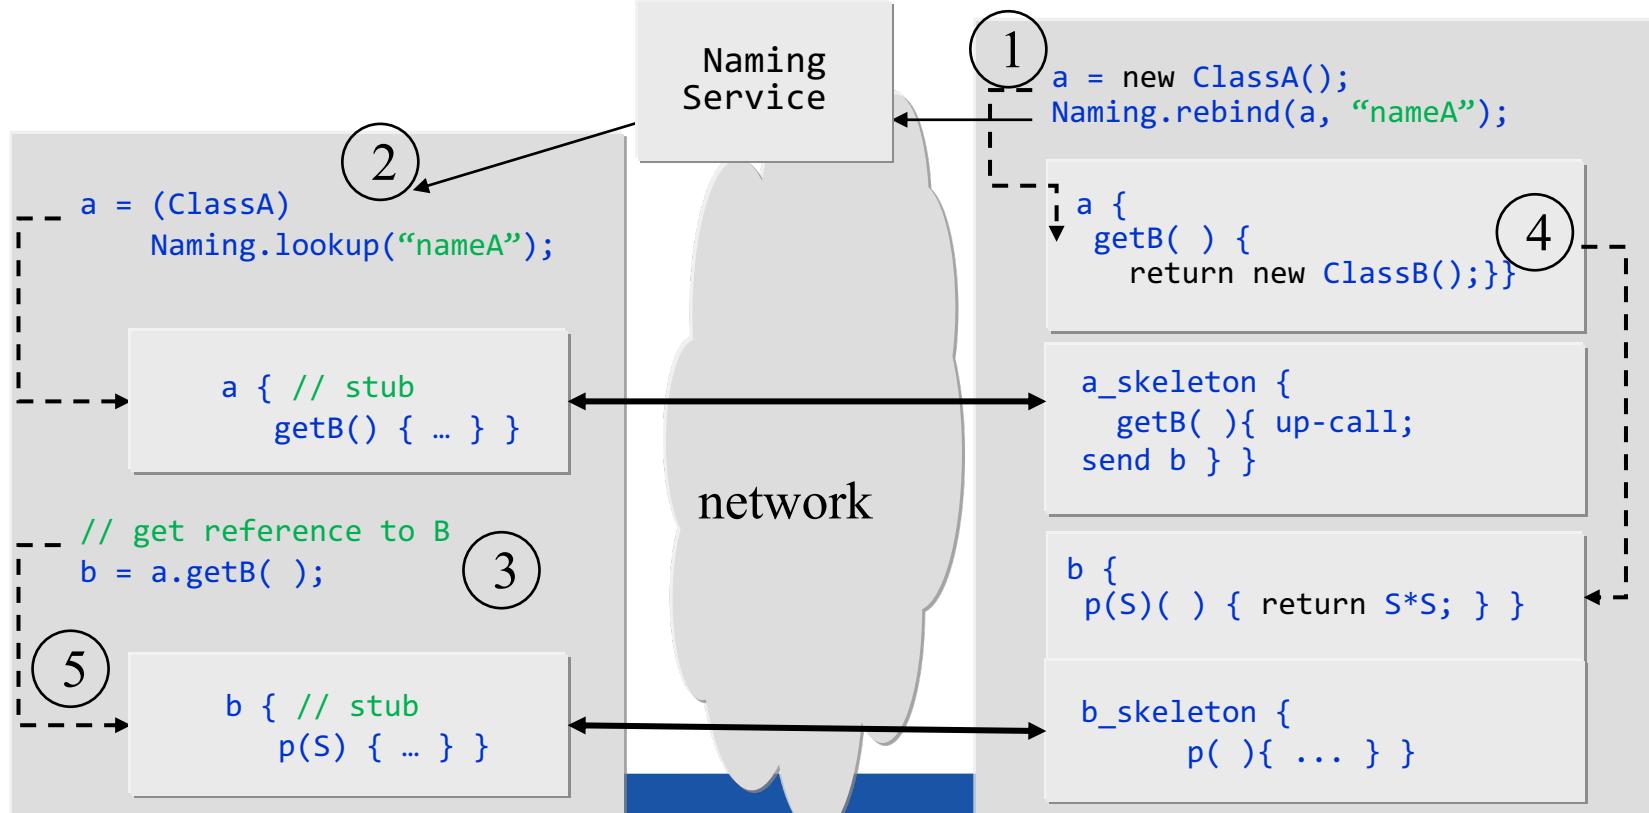
The primary reference problem: How to locate the Naming Service?

- Configuration problem: URL of the naming service

Use of Naming Service



Remote Reference in Return





Remote invocation design decisions

- failure handling: maybe / at-most-once / at-least-once
- call-by-value / call-by-reference
- message specification and encoding
- specification of resource
- procedure binder – naming service



Examples

- **SunRPC**: call-by-value, at-least-once, IDL, XDR, binder
- **JavaRMI**: call-by-value/reference, at-most-once, interface, JRMP (Java Remote Method Protocol), rmiregistry
- **Erlang**: message passing, maybe, no, ETF (External Term Format), local registry only
- **CORBA** (Common Object Request Broker Architecture): call-by-reference, IDL, ORB (Object Request Broker), tnameserv
- **Web Services**: WSDL (Web Services Description Language), UDDI (Universal Description, Discovery, and Integration)



Java RMI (Remote Method Invocation)

Java RMI is a mechanism that allows a thread in one JVM to invoke a method on an object located in another JVM.

- Provides **Java native ORB (Object Request Broker)**

The Java RMI facility allows applications running on different JVMs to interact with each other by invoking remote methods:

- Remote reference (stub) is treated as a local object.
- Method invocation on the reference causes the method to be executed on the remote JVM.
- Serialized arguments and return values are passed over network connections.
- Uses **Object streams** to pass objects “by value.”

RMI Classes and Interfaces

java.rmi.Remote

- The interface that indicates interfaces whose methods may be invoked from non-local JVM – remote interfaces.

java.rmi.Naming

- The RMI Naming Service client is used to bind a name to an object and look up an object by name.

java.rmi.RemoteException

- The common superclass for several communication-related RMI exceptions.

java.rmi.server.UnicastRemoteObject

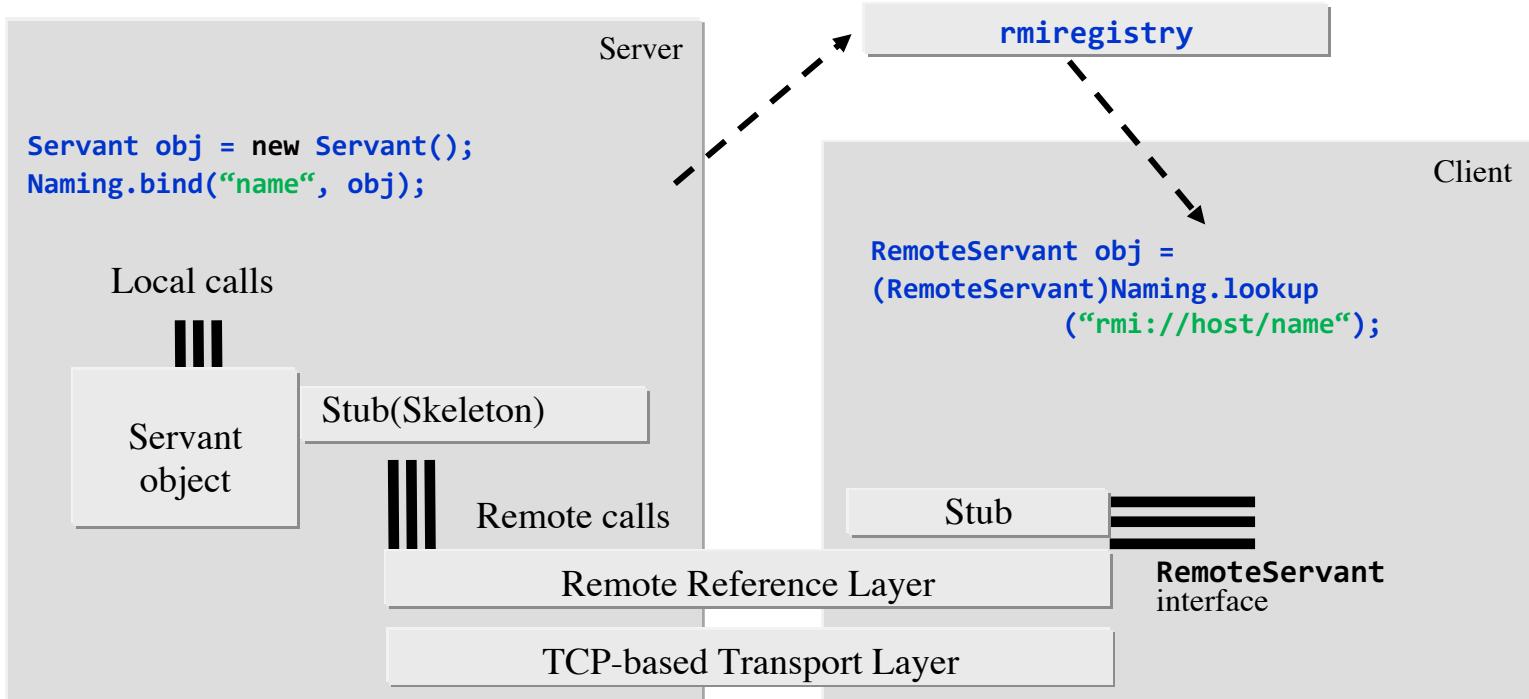
- A class that indicates a non-replicated remote object.



Developing and Executing a Distributed Application with Java RMI

1. Define a remote interface(s) that extends `java.rmi.Remote`.
2. Develop a class (a.k.a. servant class) that implements the interface.
3. Develop a server class that provides a container for servants, i.e., creates the servants and registers them at the Naming Service.
4. Develop a client class that gets a reference to a remote object(s) and calls its remote methods.
5. Compile all classes and interfaces using `javac`.
6. Start the Naming service `rmiregistry`
7. Start the server on a server host, and run the client on a client host.

Architecture of a Client-Server Application with Java RMI





Example: A Bank Manager

An application that controls accounts.

Remote interfaces:

- **Account** – deposit, withdraw, balance;
- **Bank** – create a new account, delete an account, get an account;

Classes that implement the interfaces:

- **BankImpl** – a bank servant class that implements the **Bank** interface used to create and delete accounts;
- **AccountImpl** – an account servant class implements the **Account** interface to access accounts.

Bank and Account Remote Interfaces

The **Bank** interface

```
package bankrmi;
import java.rmi.*;
import bankrmi.Account;
import bankrmi.Rejected;
public interface Bank extends Remote {
    public Account newAccount(String name) throws RemoteException, Rejected;
    public Account getAccount (String name) throws RemoteException;
    public boolean deleteAccount(String name) throws RemoteException, Rejected;
}
```

The **Account** interface

```
package bankrmi;
import java.rmi.*;
import bankrmi.Rejected;
public interface Account extends Remote {
    public float balance() throws RemoteException;
    public void deposit(float value) throws RemoteException, Rejected;
    public void withdraw(float value) throws RemoteException, Rejected;
```



A Bank Implementation

```
package bankrmi;
import java.rmi.server.UnicastRemoteObject;
import java.util.Hashtable;
import java.rmi.*;
import bankrmi.*;
public class BankImpl extends UnicastRemoteObject implements Bank {
    private String _bankname = "Noname";
    private Hashtable _accounts = new Hashtable(); // accounts
    public BankImpl(String name) throws RemoteException {
        super(); _bankname = name;
    }
    public BankImpl() throws RemoteException {
        super();
    }
    public synchronized Account newAccount(String name) throws RemoteException, Rejected {
        AccountImpl account = (AccountImpl) _accounts.get(name);
        if (account != null) {
            System.out.println("Account [" + name + "] exists!!!");
            throw new Rejected("Rejected: Bank: " + bankname + " Account for: " + name + " already exists: " + account);
        }
        account = new AccountImpl(name);
        _accounts.put(name, account);
        System.out.println("Bank: " + _bankname + " Account: " + name + " Created for " + name);
        return (Account)account;
    }
    . .
}
```



An Account Implementation

```
package bankrmi;
import java.rmi.server.UnicastRemoteObject;
import java.rmi.*;
import bankrmi.*;
public class AccountImpl extends UnicastRemoteObject implements Account {
    private float _balance = 0;
    private String _name = "noname";
    public AccountImpl(String name) throws RemoteException {
        super();
        this.name = name;
    }
    public AccountImpl() throws RemoteException {
        super();
    }
    public synchronized void deposit(float value) throws RemoteException, Rejected {
        if (value < 0) throw new Rejected("Rejected: Account " + name + ": Illegal value: " + value);
        _balance += value;
        System.out.println("Transaction: Account "+name+": deposit: $" + value + ", balance: $" + _balance);
    }
    public synchronized void withdraw(float value) throws RemoteException, Rejected {
        ...
    }
    public synchronized float balance() throws RemoteException { return _balance; }
}
```

The Server

```
package bankrmi;
import java.rmi.*;
import bankrmi.*;
public class Server {
    static final String USAGE = "java bankrmi.Server <bank_url>";
    static final String BANK = "NordBanken";
    public Server(String[] args) {
        String bankname = (args.length > 0)? args[0] : BANK;
        if (args.length > 1 || bankname.equalsIgnoreCase("-h")) {
            System.out.println(USAGE);
            System.exit(1);
        }
        try {
            Bank bankobj = (Bank)(new BankImpl(bankname));
            Naming.rebind(bankname, bankobj);
            System.out.println(bankobj + " is ready.");
        } catch (Exception e) { System.out.println(e); }
        Object sync = new Object();
        synchronized(sync) { try { sync.wait();} catch (Exception ie) {}}
    }
    public static void main(String[] args) {
        new Server(args).start();
    }
}
```



A Client

```
package bankrmi;
import bankrmi.*;
import java.rmi.*;
public class SClient {
    static final String USAGE = "java Client <bank_url> <client> <value>";
    String bankname = "Noname", clientname = "Noname"; // defaults
    float value = 100;
    public SClient(String[] args) {
        // Read and parse command line arguments (see Usage above)
        ...
        try {
            Bank bankobj = (Bank)Naming.lookup(bankname);
            Account account = bankobj.newAccount(clientname);
            account.deposit(value);
            System.out.println (clientname + "'s account: $" + account.balance());
        } catch (Exception e) {
            System.out.println("The runtime failed: " + e);
            System.exit(0);
        }
    }
    public static void main(String[] args) {
        new SClient(args);
    }
}
```



Summary

Implementations of remote invocations: procedures, methods,
messages to processes,
have fundamental problems that need to be solved.

Try to see similarities between different implementations.

When they differ, is it fundamentally different or just
implementation details?